

EVIDENCE MAPPED STUDY EDITION

MANIPULATION

The Fool and the Wise

INFLUENCE / DECEPTION / POWER / AI



J. NASR

Why this book exists

Influence becomes easier to examine when it has a name, an observable mechanism and a traceable source. This book joins those elements across everyday conversation, organizations, commerce, relationships, interfaces, information systems and AI. Its purpose is to improve the reader's ability to explain what is happening, distinguish evidence from suspicion and design more accountable applications.

The fool and the wise are positions we can move between, not permanent types of people. The unexamined conclusion is the point of departure. The studied conclusion asks what was said, what was inferred, which alternatives were available and what could change the interpretation. Being deceived is not proof of low intelligence or moral failure.

This edition retains the full mechanism inventory of the supplied research atlas: 66 Class 1 implementation entries and 144 Class 0 entries, plus 13 separately identified hypotheses. The source assigns E to 97 entries and C to 113. These counts describe the catalogue, not independent psychological discoveries. Related entries overlap by design.

The AI layer contains three mapped directions per entry, or 669 directions in total, and a portfolio of 40 cross-mechanism product concepts. Many directions reuse a common architecture. They should not be advertised as 669 unique inventions or tested businesses. The companion register preserves every original application mapping.

Evidence and implementation are different questions

E means the source atlas identifies a directly documented mechanism or an authoritative legal or technical category. It does not mean strong evidence of universal effectiveness. A regulatory definition and a behavioral experiment establish different kinds of facts. E is retained as a source classification; this edition has not independently revalidated every exact row-to-source assignment.

C means the specific label or configuration is an editorial synthesis assembled from documented components. Its name is not presented as canonical academic terminology. H means the precise configuration remains a hypothesis within this source map. H does not establish novelty or imply that the pattern has never occurred.

Class 1 identifies an implementation that can be legitimate under the conditions described. It does not make every use harmless. Class 0 groups deceptive, coercive, exploitative or adversarial configurations for analysis. The notation is project metadata, not an assessment of a person's character, a universal legal conclusion or a permission system.

A true claim, a real identity and a visible alternative are useful considerations, but no single checklist settles every case. A person can face genuine information and still have very limited practical freedom. Conversely, an awkward interface can reflect poor design without an intention to manipulate.

Truthful words and false impressions

The sentence can be true while the conclusion is false

Communication conveys more than literal assertions. Listeners infer relevance, completeness and the reason a fact was selected. Research on paltering examines the active use of truthful statements to mislead; work on communicative norms also shows how misleading impressions can arise without the same intent. The distinction matters because a misleading outcome alone does not establish deliberate deception. [R42, R43]

Consider three levels of analysis. First, record the exact statement. Second, describe the conclusion the audience is likely to draw in that setting. Third, investigate whether the speaker knew material context that undermines that conclusion. The first level concerns text, the second interpretation and the third evidence about intent. Collapsing them creates false accusations as easily as it misses calculated ambiguity.

The composite labels in this chapter isolate different parts of that process. Context laundering changes the setting of a true statement. Deceptive sequence manipulation changes when corrective information becomes available. Inference laundering concerns the later defense that the false conclusion was never explicitly spoken. These are study distinctions derived from documented components, and several can describe the same incident.

A life scenario

A studio is asked whether an independent auditor approved its security. Its representative replies that the system uses tools also used by major banks. That may be true, but it does not answer the audit question. The listener must avoid converting a statement about tools into evidence of independent approval.

Application routes for this chapter

The applications should highlight the sentence, the possible inference and the missing verification. A model cannot determine private knowledge or intent from polished wording alone.

Route A provides a personal evidence review. Route B supports a professional audit and verification trail. Route C creates labelled teaching and evaluation cases. Each entry below gives a specific proposed output; the companion register retains all three original mappings.

All detection directions are proposals. Evidence, competing explanations, privacy and a human review step matter more than an unsupported severity score.

ATLAS 05 / 13 ENTRIES

05.01 Lying by commission

E / CLASS 0 / Deception

Uses a knowingly false assertion to induce a false belief.

Life example. A seller knowingly says a used car has never been in an accident when records show that it has.

AI application focus. A claim-verification sheet comparing an assertion with available records while keeping knowledge and intent unresolved.

Source anchors: [R42]

05.02 Lying by omission

E / CLASS 0 / Deception

Passively withholds relevant information in a context where the omission creates a misleading impression.

Life example. A recruiter knows a role is temporary but omits that fact while presenting it as stable long-term employment.

AI application focus. A missing-information checklist tied to the specific question or expectation created by the conversation.

Source anchors: [R42]

05.03 Paltering

E / CLASS 0 / Deception/negotiation

Actively uses truthful statements to create a misleading impression.

Life example. A negotiator truthfully says "sales increased 20%" while intentionally omitting that profit and customer count collapsed, leading the listener to infer overall business improvement.

AI application focus. A comparison of literal statements, likely impressions and the context needed to test those impressions.

Source anchors: [R42, R43]

05.04 Omission engineering

C / CLASS 0 / Communication

Strategically withholds material context while surrounding claims remain literally true.

Life example. A property advertisement lists attractive features but strategically omits a known mandatory building charge that materially changes the deal.

AI application focus. A material-context review showing which omitted facts could change the decision.

Source anchors: [R42, R43]

05.05 Inference baiting

C / CLASS 0 / Communication

Arranges true or ambiguous cues so the audience is expected to infer a consequential falsehood that is never asserted.

Life example. A spokesperson mentions several true details that strongly suggest a project was independently audited even though no independent audit occurred.

AI application focus. An inference map connecting supplied cues with a consequential conclusion that has not been established.

Source anchors: [R42, R43]

05.06 Listener-completed deception

C / CLASS 0 / Communication

Relies on the target to supply the final false inference from technically true inputs.

Life example. A person says only true things about being "at work" so their partner concludes they were in the office, while they intentionally rely on that unstated conclusion.

AI application focus. A transcript annotation separating spoken facts from conclusions supplied by the listener.

Source anchors: [R42, R43]

05.07 Inference laundering

C / CLASS 0 / Communication

Induces a false impression and later defends it by pointing out that the false conclusion was never explicitly stated.

Life example. After creating a misleading impression, a salesperson defends themselves with "I never actually said that was included."

AI application focus. A timeline connecting the original impression with later defenses about what was never explicitly said.

Source anchors: [R42, R43]

05.08 Plausible-deniability persuasion

C / CLASS 0 / Communication

Designs a misleading influence attempt so intent can be denied because individual statements remain literally defensible.

Life example. A manager wordsmiths a claim so it implies guaranteed promotion while retaining enough ambiguity to deny making a promise later.

AI application focus. A wording review identifying ambiguities that leave a consequential promise unclear.

Source anchors: [R42, R43]

05.09 Ambiguity weaponization

C / CLASS 0 / Communication

Preserves ambiguity because the target is likely to resolve it in the influencer's preferred, materially false direction.

Life example. A contract discussion deliberately uses "may" and "typically" where the listener is expected to hear a stronger promise than the words support.

AI application focus. A list of reasonable interpretations and a precise clarification question for each.

Source anchors: [R43]

05.10 Context laundering

C / CLASS 0 / Communication

Moves a true statement into a context where normal assumptions materially alter its apparent meaning.

Life example. A genuine quote about one limited situation is placed beside unrelated imagery so viewers infer that it describes the broader event.

AI application focus. A source-context comparison showing what changes when a statement is moved into a different setting.

Source anchors: [R43]

05.11 Deceptive sequence manipulation

C / CLASS 0 / Communication

Orders information so early material fixes an interpretation while decisive context is delayed or buried.

Life example. A presentation leads with impressive percentages and delays the small sample size until the final slide, after the audience has formed its conclusion.

AI application focus. A presentation timeline showing when decisive information becomes available relative to the main claim.

Source anchors: [R43]

05.12 Semantic threat inflation

C / CLASS 0 / Communication

Uses exaggerated threat language that materially overstates actual risk.

Life example. A campaign describes a modest operational problem as a "catastrophic threat" despite evidence showing limited risk.

AI application focus. A comparison of threat language with the magnitude and uncertainty of the documented risk.

Source anchors: [R09, R38]

05.13 Material misleading implication

C / CLASS 0 / Consumer/deception

Creates a false net impression through wording, format, or omission even when some individual claims are true.

Life example. An advertisement says "clinically studied ingredients" in a way that leaves consumers with the false impression that the finished product itself was clinically proven.

AI application focus. A review of the overall impression against what the available evidence actually establishes.

Source anchors: [R21, R42]

Study question

What additional fact would resolve the ambiguity without speculating about character?

Model answer. Whether an independent audit occurred, its scope and its result. Request that evidence directly rather than debating how impressive the surrounding statement sounds.

Transfer exercise. Find a different situation that uses one mechanism in this chapter. Describe one fact that would support your classification and one that would weaken it.

Numbers images and evidence

A correct number can carry an incorrect scale

A number answers a question only when the denominator, comparison, time period and measurement are clear. Relative change describes a relationship between values; absolute change supplies magnitude. A visual encoding adds another layer because lengths, areas and axes influence the apparent difference. Studies of misleading graphs caution against assuming identical effects across all tasks and contexts. [R44, R45, R54]

Some distortions concern the data and others concern what the data are said to represent. Counting views is not intrinsically wrong. Calling those views supporters replaces one construct with another. Similarly, a genuine citation can fail to support the stronger sentence written beside it. Evidence quality therefore requires a chain from claim to measurement to source, not a count of impressive-looking references.

Source independence is especially easy to overstate. Five articles that reproduce one press release provide five distribution points but not five independent observations. This book itself requires the same scrutiny: repeated reference IDs and application-layer duplicates must not be counted as separate corroborating studies. A full URL helps locate a source; it does not by itself establish that the associated interpretation is correct.

A life scenario

An ad reports a 100 percent rise in complaints. The actual counts are one complaint last month and two this month, while transactions increased from one hundred to one thousand. Complaint count rose; complaint rate fell from 1 percent to 0.2 percent. Both calculations are correct, but they answer different questions.

Application routes for this chapter

Apps should return the original values, recomputed quantities and source trail. Numerical checking belongs in reproducible calculations; language generation can explain those calculations.

Route A provides a personal evidence review. Route B supports a professional audit and verification trail. Route C creates labelled teaching and evaluation cases. Each entry below gives a specific proposed output; the companion register retains all three original mappings.

All detection directions are proposals. Evidence, competing explanations, privacy and a human review step matter more than an unsupported severity score.

ATLAS 06 / 12 ENTRIES

06.01 Relative-risk-only amplification

E / CLASS 0 / Risk communication

Presents only relative change when hidden absolute risks materially exaggerate perceived magnitude.

Life example. A drug story reports that risk "doubled" from 1 in 10,000 to 2 in 10,000 but never states the absolute numbers.

AI application focus. An absolute-risk calculation beside the relative headline, with the starting and ending values visible.

Source anchors: [R54]

06.02 Denominator concealment

C / CLASS 0 / Quantitative communication

Uses percentages or frequencies without the denominator needed to judge scale.

Life example. A report says "50% of users complained" without revealing that only two users were surveyed.

AI application focus. A statistic annotated with its denominator, sample size and collection period.

Source anchors: [R54]

06.03 Base-rate hiding

C / CLASS 0 / Quantitative communication

Emphasizes conditional or dramatic cases while omitting the relevant population baseline.

Life example. A news graphic highlights ten dramatic failures but omits that they occurred among millions of successful transactions.

AI application focus. A claim review that requests the relevant population baseline before interpreting dramatic cases.

Source anchors: [R54]

06.04 Truncated-axis visual distortion

E / CLASS 0 / Data visualization

Uses axis truncation to exaggerate apparent differences.

Life example. A bar chart starts its vertical axis at 95 instead of zero so a change from 98 to 100 looks enormous.

AI application focus. A chart comparison using the original and alternative axes, with the actual numerical difference stated.

Source anchors: [R44, R45]

06.05 Visual-scale distortion

E / CLASS 0 / Data visualization

Changes graphical scale or proportional representation so visual magnitude misstates numerical magnitude.

Life example. Two pie-like graphics use different circle areas that are not proportional to the underlying percentages, making one category appear much larger.

AI application focus. A proportionality check comparing the graphic's visual encoding with the underlying values.

Source anchors: [R44, R45]

06.06 Temporal cherry-picking

C / CLASS 0 / Evidence presentation

Selects a time window that manufactures a trend impression inconsistent with the relevant period.

Life example. A company promotes a chart beginning immediately after an unusually bad month so the selected window creates the appearance of rapid sustained growth.

AI application focus. A trend comparison across the claimed period and a justified wider window.

Source anchors: [R38, R45]

06.07 Reference-class distortion

C / CLASS 0 / Evidence presentation

Chooses an artificially weak, irrelevant, or non-comparable comparator to manufacture superiority.

Life example. A school advertises that its graduates earn more than a comparison group drawn from a much poorer region rather than comparable schools.

AI application focus. A comparison-group review identifying material differences that could explain apparent superiority.

Source anchors: [R09, R54]

06.08 Metric laundering

C / CLASS 0 / Evidence presentation

Presents a proxy metric as if it measured a stronger construct, such as treating impressions as approval or registrations as customers.

Life example. A campaign reports millions of social-media impressions as though they were millions of supporters.

AI application focus. A metric dictionary separating what was counted from the stronger outcome claimed.

Source anchors: [R38]

06.09 Evidence cropping

C / CLASS 0 / Evidence presentation

Removes material surrounding context so authentic evidence takes on a substantially different meaning.

Life example. A video shares ten seconds of a real interview while removing the preceding question that materially changes what the answer meant.

AI application focus. A side-by-side comparison of the cropped extract and its material surrounding context.

Source anchors: [R38, R43]

06.10 Quote laundering

C / CLASS 0 / Information integrity

Repetition turns a fragment, paraphrase, or uncertain attribution into an apparently verified quotation.

Life example. A saying is repeatedly copied across quote sites until later pages confidently attribute it to a famous author despite no primary source.

AI application focus. An attribution trail ending in a primary text or an explicit unresolved attribution.

Source anchors: [R38, R40]

06.11 Citation/source laundering

C / CLASS 0 / Information integrity

Routes weak, circular, or derivative claims through citations so they appear independently corroborated.

Life example. Dozens of blogs cite one another for the same claim, making a single unsupported origin look like many independent sources.

AI application focus. A citation graph showing direct evidence, derivative repetition and unsupported origins.

Source anchors: [R38, R40]

06.12 Consensus laundering

C / CLASS 0 / Information integrity

Presents non-independent sources or repeated copies as independent agreement.

Life example. A presentation counts five websites repeating the same press release as five independent confirmations.

AI application focus. A count of independent evidential origins rather than a count of repeating webpages.

Source anchors: [R38, R39]

Study question

Which claim would you make if the question concerned service quality?

Model answer. Report the counts, transaction denominators, rates and period. Then investigate comparability and collection methods before attributing the difference to quality.

Transfer exercise. Find a different situation that uses one mechanism in this chapter. Describe one fact that would support your classification and one that would weaken it.

Continue the study

The complete edition contains fifteen core chapters, 210 mechanism entries, thirteen hypotheses, six extended case studies, forty AI product concepts and an alphabetical index. A separate workbook and complete application register accompany the book.

This sample reproduces selected pages, so the original book page numbers are retained. The examples are fictional teaching illustrations. E and C are source evidence classifications, not certainty scores.

The full edition is prepared for release. Purchase details will appear on the author's book page once checkout opens.

References used in this excerpt

R01 Noggle, R. The Ethics of Manipulation. Stanford Encyclopedia of Philosophy. Substantive revision 31 March 2026.

<https://plato.stanford.edu/entries/ethics-manipulation/>

R09 Kuhberger, A. (1998). The Influence of Framing on Risky Decisions: A Meta-analysis. *Organizational Behavior and Human Decision Processes*, 75(1), 23-55.

<https://doi.org/10.1006/obhd.1998.2781>

R21 Federal Trade Commission. Native Advertising: A Guide for Businesses.

<https://www.ftc.gov/business-guidance/resources/native-advertising-guide-businesses>

R38 van der Linden, S. (2022). Misinformation: susceptibility, spread, and interventions to immunize the public. *Nature Medicine*, 28, 460-467.

<https://doi.org/10.1038/s41591-022-01713-6>

R39 Traberg, C. S., Harjani, T., Roozenbeek, J., et al. (2024). The persuasive effects of social cues and source effects on misinformation susceptibility. *Scientific Reports*, 14, 4205.

<https://doi.org/10.1038/s41598-024-54030-y>

R40 Edelson, S. M., Reyna, V. F., Singh, A., & Roue, J. E. (2024). The Psychology of Misinformation Across the Lifespan. *Annual Review of Developmental Psychology*, 6, 425-454.

<https://doi.org/10.1146/annurev-devpsych-010923-093547>

R42 Rogers, T., Zeckhauser, R., Gino, F., Norton, M. I., & Schweitzer, M. E. (2017). Artful paltering: The risks and rewards of using truthful statements to mislead others. *Journal of Personality and Social Psychology*, 112(3), 456-473.

<https://doi.org/10.1037/pspi0000081>

R43 Powell, D., Bian, L., & Markman, E. M. (2020). When intents to educate can misinform: Inadvertent paltering through violations of communicative norms. *PLOS ONE*, 15(5), e0230360.

<https://doi.org/10.1371/journal.pone.0230360>

R44 Vargas Restrepo, C., Yang, B. W., Stanley, M. L., & Marsh, E. J. (2021). Truncating Bar Graphs Persistently Misleads Viewers. *Journal of Applied Research in Memory and Cognition*, 10(2), 298-311.

<https://doi.org/10.1016/j.jarmac.2020.10.002>

R45 Driessen, J. E. P., Vos, D. A. C., Smeets, I., & Albers, C. J. (2022). Misleading graphs in context: Less misleading than expected. *PLOS ONE*, 17, e0265823.

<https://doi.org/10.1371/journal.pone.0265823>

R54 Noordzij, M., van Diepen, M., Caskey, F. C., & Jager, K. J. (2017). Relative risk versus absolute risk: one cannot be interpreted without the other. *Nephrology Dialysis Transplantation*, 32(suppl_2), ii13-ii18.

<https://doi.org/10.1093/ndt/gfw465>